



**Министерство науки и высшего образования Российской
Федерации
Федеральное государственное бюджетное образовательное
учреждение
высшего образования
«Московский государственный технический университет
имени Н.Э. Баумана
(национальный исследовательский университет)»
(МГТУ им. Н.Э. Баумана)**

ФАКУЛЬТЕТ ИНФОРМАТИКА И СИСТЕМЫ УПРАВЛЕНИЯ

КАФЕДРА СИСТЕМЫ ОБРАБОТКИ ИНФОРМАЦИИ И УПРАВЛЕНИЯ

Отчёт к лабораторным работам по курсу

«Методы машинного обучения»

**Лабораторная работа №6 «Обучение на основе глубоких Q-
сетей»**

Выполнил:

студент(ка) группы ИУ5И-21М Лю Бэйбэй

подпись, дата

Проверил:

к.т.н., доц., Виноградовой М.В.

подпись, дата

Москва, 2022 г.

1. описание задания

На основе рассмотренных на лекции примеров реализуйте алгоритм DQN.

В качестве среды можно использовать классические среды (в этом случае используется полносвязная архитектура нейронной сети).

В качестве среды можно использовать игры Atari (в этом случае используется сверточная архитектура нейронной сети).

В случае реализации среды на основе сверточной архитектуры нейронной сети
+1 балл за экзамен.

2. Текст программы и экранные формы с примерами выполнения программы.

```
class DeepQNetwork:
    def __init__(
        self,
        n_actions,
        n_features,
        learning_rate=0.01,
        reward_decay=0.9,
        e_greedy=0.9,
        replace_target_iter=300,
        memory_size=500,
        batch_size=32,
        e_greedy_increment=None,
        output_graph=False,
    ):
        self.n_actions = n_actions
        self.n_features = n_features
        self.lr = learning_rate
        self.gamma = reward_decay
        self.epsilon_max = e_greedy
        self.replace_target_iter = replace_target_iter
        self.memory_size = memory_size
        self.batch_size = batch_size
        self.epsilon_increment = e_greedy_increment
        self.epsilon = 0 if e_greedy_increment is not None else
self.epsilon_max

        # total learning step
        self.learn_step_counter = 0

        # initialize zero memory [s, a, r, s_]
        self.memory = np.zeros((self.memory_size, n_features * 2 + 2))

        # consist of [target_net, evaluate_net]
        self._build_net()
        t_params = tf.get_collection('target_net_params')
        e_params = tf.get_collection('eval_net_params')
        self.replace_target_op = [tf.assign(t, e) for t, e in zip(t_params,
e_params)]
```

```

self.sess = tf.Session()

if output_graph:
    # $ tensorboard --logdir=logs
    # tf.train.SummaryWriter soon be deprecated, use following
    tf.summary.FileWriter("logs/", self.sess.graph)

self.sess.run(tf.global_variables_initializer())
self.cost_hist = []

def _build_net(self):
    # ----- build evaluate_net -----
    self.s = tf.placeholder(tf.float32, [None, self.n_features], name='s') #
input
    self.q_target = tf.placeholder(tf.float32, [None, self.n_actions],
name='Q_target') # for calculating loss
    with tf.variable_scope('eval_net'):
        # c_names(collections_names) are the collections to store
variables
        c_names, n_l1, w_initializer, b_initializer = \
            ['eval_net_params', tf.GraphKeys.GLOBAL_VARIABLES],
10, \
            tf.random_normal_initializer(0., 0.3),
tf.constant_initializer(0.1) # config of layers

        # first layer. collections is used later when assign to target net
        with tf.variable_scope('l1'):
            w1 = tf.get_variable('w1', [self.n_features, n_l1],
initializer=w_initializer, collections=c_names)
            b1 = tf.get_variable('b1', [1, n_l1], initializer=b_initializer,
collections=c_names)
            l1 = tf.nn.relu(tf.matmul(self.s, w1) + b1)

        # second layer. collections is used later when assign to target net
        with tf.variable_scope('l2'):
            w2 = tf.get_variable('w2', [n_l1, self.n_actions],
initializer=w_initializer, collections=c_names)
            b2 = tf.get_variable('b2', [1, self.n_actions],
initializer=b_initializer, collections=c_names)
            self.q_eval = tf.matmul(l1, w2) + b2

    with tf.variable_scope('loss'):

```

```

        self.loss = tf.reduce_mean(tf.squared_difference(self.q_target,
self.q_eval))
        with tf.variable_scope('train'):
            self._train_op =
tf.train.RMSPropOptimizer(self.lr).minimize(self.loss)

        # ----- build target_net -----
        self.s_ = tf.placeholder(tf.float32, [None, self.n_features], name='s_')
# input
        with tf.variable_scope('target_net'):
            # c_names(collections_names) are the collections to store
variables
            c_names = ['target_net_params',
tf.GraphKeys.GLOBAL_VARIABLES]

            # first layer. collections is used later when assign to target net
            with tf.variable_scope('l1'):
                w1 = tf.get_variable('w1', [self.n_features, n_l1],
initializer=w_initializer, collections=c_names)
                b1 = tf.get_variable('b1', [1, n_l1], initializer=b_initializer,
collections=c_names)
                l1 = tf.nn.relu(tf.matmul(self.s_, w1) + b1)

            # second layer. collections is used later when assign to target net
            with tf.variable_scope('l2'):
                w2 = tf.get_variable('w2', [n_l1, self.n_actions],
initializer=w_initializer, collections=c_names)
                b2 = tf.get_variable('b2', [1, self.n_actions],
initializer=b_initializer, collections=c_names)
                self.q_next = tf.matmul(l1, w2) + b2

def store_transition(self, s, a, r, s_):
    if not hasattr(self, 'memory_counter'):
        self.memory_counter = 0

    transition = np.hstack((s, [a, r], s_))

    # replace the old memory with new memory
    index = self.memory_counter % self.memory_size
    self.memory[index, :] = transition

    self.memory_counter += 1

def choose_action(self, observation):

```

```

# to have batch dimension when feed into tf placeholder
observation = observation[np.newaxis, :]

if np.random.uniform() < self.epsilon:
    # forward feed the observation and get q value for every actions
    actions_value = self.sess.run(self.q_eval, feed_dict={self.s:
observation})
    action = np.argmax(actions_value)
else:
    action = np.random.randint(0, self.n_actions)
return action

def learn(self):
    # check to replace target parameters
    if self.learn_step_counter % self.replace_target_iter == 0:
        self.sess.run(self.replace_target_op)
        print('\ntarget_params_replaced\n')

    # sample batch memory from all memory
    if self.memory_counter > self.memory_size:
        sample_index = np.random.choice(self.memory_size,
size=self.batch_size)
    else:
        sample_index = np.random.choice(self.memory_counter,
size=self.batch_size)
    batch_memory = self.memory[sample_index, :]

    q_next, q_eval = self.sess.run(
        [self.q_next, self.q_eval],
        feed_dict={
            self.s_ : batch_memory[:, -self.n_features:], # fixed params
            self.s: batch_memory[:, :self.n_features], # newest params
        })

    # change q_target w.r.t q_eval's action
    q_target = q_eval.copy()

    batch_index = np.arange(self.batch_size, dtype=np.int32)
    eval_act_index = batch_memory[:, self.n_features:].astype(int)
    reward = batch_memory[:, self.n_features + 1]

    q_target[batch_index, eval_act_index] = reward + self.gamma *
np.max(q_next, axis=1)

```

```

        # train eval network
        _, self.cost = self.sess.run([self._train_op, self.loss],
                                      feed_dict={self.s:
batch_memory[:, :self.n_features],
                                                self.q_target:
q_target})

        self.cost_his.append(self.cost)

        # increasing epsilon
        self.epsilon = self.epsilon + self.epsilon_increment if self.epsilon <
self.epsilon_max else self.epsilon_max
        self.learn_step_counter += 1

    def plot_cost(self):
        import matplotlib.pyplot as plt
        plt.plot(np.arange(len(self.cost_his)), self.cost_his)
        plt.ylabel('Cost')
        plt.xlabel('training steps')
        plt.show()

UNIT = 40    # pixels
MAZE_H = 4   # grid height
MAZE_W = 4   # grid width

class Maze(tk.Tk, object):
    def __init__(self):
        super(Maze, self).__init__()
        self.action_space = ['u', 'd', 'l', 'r']
        self.n_actions = len(self.action_space)
        self.n_features = 2
        self.title('maze')
        self.geometry('{0}x{1}'.format(MAZE_W * UNIT, MAZE_H *
UNIT))

        self._build_maze()

    def _build_maze(self):
        self.canvas = tk.Canvas(self, bg='white',
                                height=MAZE_H * UNIT,
                                width=MAZE_W * UNIT)

        # create grids
        for c in range(0, MAZE_W * UNIT, UNIT):
            x0, y0, x1, y1 = c, 0, c, MAZE_H * UNIT

```

```

        self.canvas.create_line(x0, y0, x1, y1)
    for r in range(0, MAZE_H * UNIT, UNIT):
        x0, y0, x1, y1 = 0, r, MAZE_W * UNIT, r
        self.canvas.create_line(x0, y0, x1, y1)

    # create origin
    origin = np.array([20, 20])

    # hell
    hell1_center = origin + np.array([UNIT * 2, UNIT])
    self.hell1 = self.canvas.create_rectangle(
        hell1_center[0] - 15, hell1_center[1] - 15,
        hell1_center[0] + 15, hell1_center[1] + 15,
        fill='black')
    # hell
    # hell2_center = origin + np.array([UNIT, UNIT * 2])
    # self.hell2 = self.canvas.create_rectangle(
    #     hell2_center[0] - 15, hell2_center[1] - 15,
    #     hell2_center[0] + 15, hell2_center[1] + 15,
    #     fill='black')

    # create oval
    oval_center = origin + UNIT * 2
    self.oval = self.canvas.create_oval(
        oval_center[0] - 15, oval_center[1] - 15,
        oval_center[0] + 15, oval_center[1] + 15,
        fill='yellow')

    # create red rect
    self.rect = self.canvas.create_rectangle(
        origin[0] - 15, origin[1] - 15,
        origin[0] + 15, origin[1] + 15,
        fill='red')

    # pack all
    self.canvas.pack()

    def reset(self):
        self.update()
        time.sleep(0.1)
        self.canvas.delete(self.rect)
        origin = np.array([20, 20])
        self.rect = self.canvas.create_rectangle(
            origin[0] - 15, origin[1] - 15,

```



```

        origin[0] + 15, origin[1] + 15,
        fill='red')
    # return observation
    return (np.array(self.canvas.coords(self.rect)[:2]) -
np.array(self.canvas.coords(self.oval)[:2]))/(MAZE_H*UNIT)

def step(self, action):
    s = self.canvas.coords(self.rect)
    base_action = np.array([0, 0])
    if action == 0:    # up
        if s[1] > UNIT:
            base_action[1] -= UNIT
    elif action == 1:    # down
        if s[1] < (MAZE_H - 1) * UNIT:
            base_action[1] += UNIT
    elif action == 2:    # right
        if s[0] < (MAZE_W - 1) * UNIT:
            base_action[0] += UNIT
    elif action == 3:    # left
        if s[0] > UNIT:
            base_action[0] -= UNIT

    self.canvas.move(self.rect, base_action[0], base_action[1])    # move
agent

    next_coords = self.canvas.coords(self.rect)    # next state

    # reward function
    if next_coords == self.canvas.coords(self.oval):
        reward = 1
        done = True
    elif next_coords in [self.canvas.coords(self.hell1)]:
        reward = -1
        done = True
    else:
        reward = 0
        done = False
    s_ = (np.array(next_coords[:2]) -
np.array(self.canvas.coords(self.oval)[:2]))/(MAZE_H*UNIT)
    return s_, reward, done

def render(self):
    # time.sleep(0.01)
    self.update()

```

```

def run_maze():
    step = 0
    for episode in range(300):
        # initial observation
        observation = env.reset()

        while True:
            # fresh env
            env.render()

            # RL choose action based on observation
            action = RL.choose_action(observation)

            # RL take action and get next observation and reward
            observation_, reward, done = env.step(action)

            RL.store_transition(observation, action, reward, observation_)

            if (step > 200) and (step % 5 == 0):
                RL.learn()

            # swap observation
            observation = observation_

            # break while loop when end of this episode
            if done:
                break
            step += 1

        # end of game
        print('game over')
        env.destroy()

if __name__ == "__main__":
    # maze game
    env = Maze()
    RL = DeepQNetwork(env.n_actions, env.n_features,
                      learning_rate=0.01,
                      reward_decay=0.9,
                      e_greedy=0.9,
                      replace_target_iter=200,
                      memory_size=2000,
                      # output_graph=True

```

```
)  
env.after(100, run_maze)  
env.mainloop()  
RL.plot_cost()
```

